PyTorch live talks

Capsules and routing techniques

Rationale

- CNN *can* deal with *translation*
- For other *affine transformations*
  - Exponential replicas of spatial feature detectors
  - Exponentially more labelled data

- Solution
  - Efficiently encode *viewpoint invariant knowledge*

Capsules

- Groups of neurons characterising an entity in the image
- Properties include
  - Pose, deformation, velocity, albedo, hue, texture…
- Role
  - Invert the rendering process
  - 2D camera projection ⟶ 3D abstract model

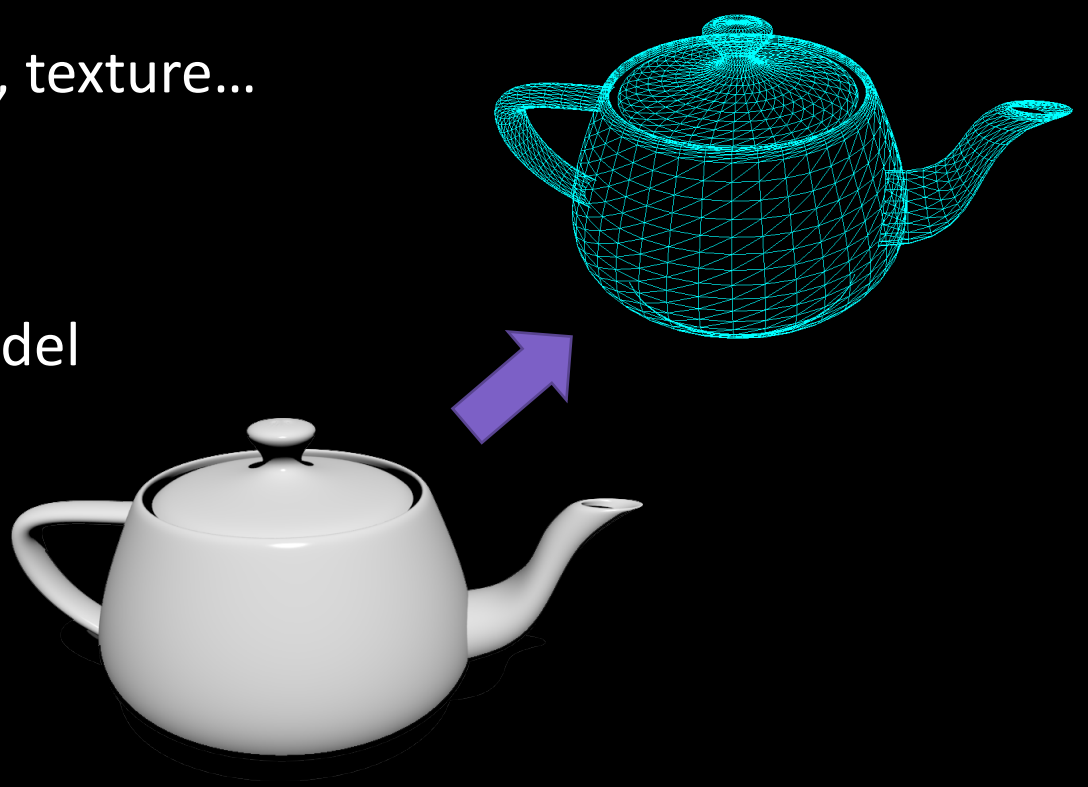

Capsules

$$\boldsymbol{v}_j = \frac{\|\boldsymbol{s}_j\|^2}{1 + \|\boldsymbol{s}_j\|^2} \frac{\boldsymbol{s}_j}{\|\boldsymbol{s}_j\|}$$

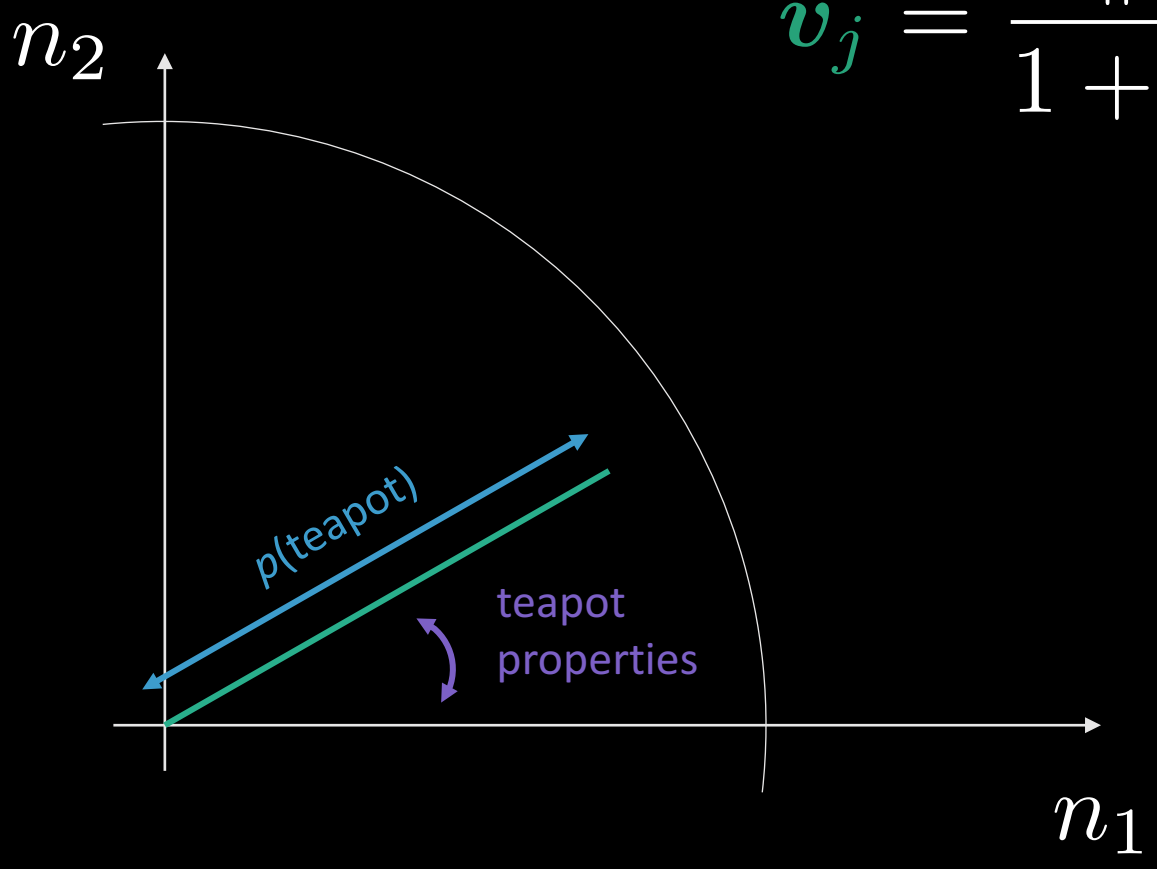

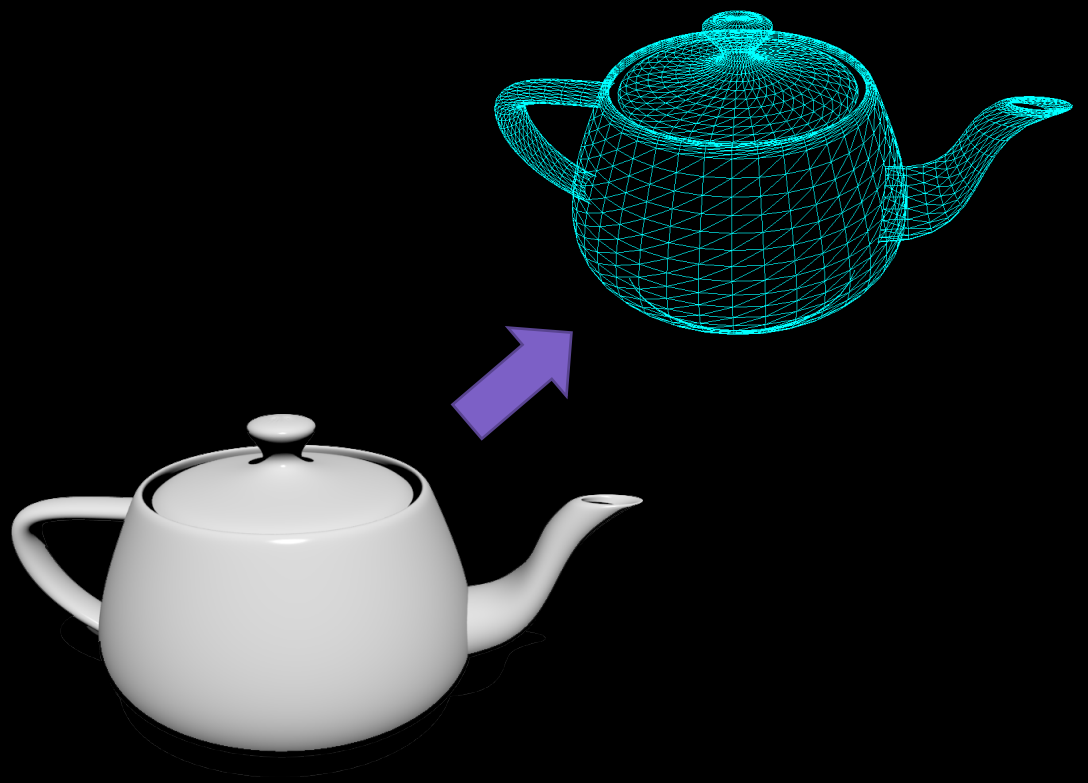

Capsules

$$\boldsymbol{v}_j = \frac{\|\boldsymbol{s}_j\|^2}{1 + \|\boldsymbol{s}_j\|^2} \frac{\boldsymbol{s}_j}{\|\boldsymbol{s}_j\|}$$

Capsule output

$$\boldsymbol{s}_j = \sum_i c_{ij} \hat{\boldsymbol{u}}_{j|i}$$

Capsule input

$$\hat{\boldsymbol{u}}_{j|i} = \boldsymbol{W}_{ij} \boldsymbol{u}_i$$

Prediction vectors

learnt by back-prop

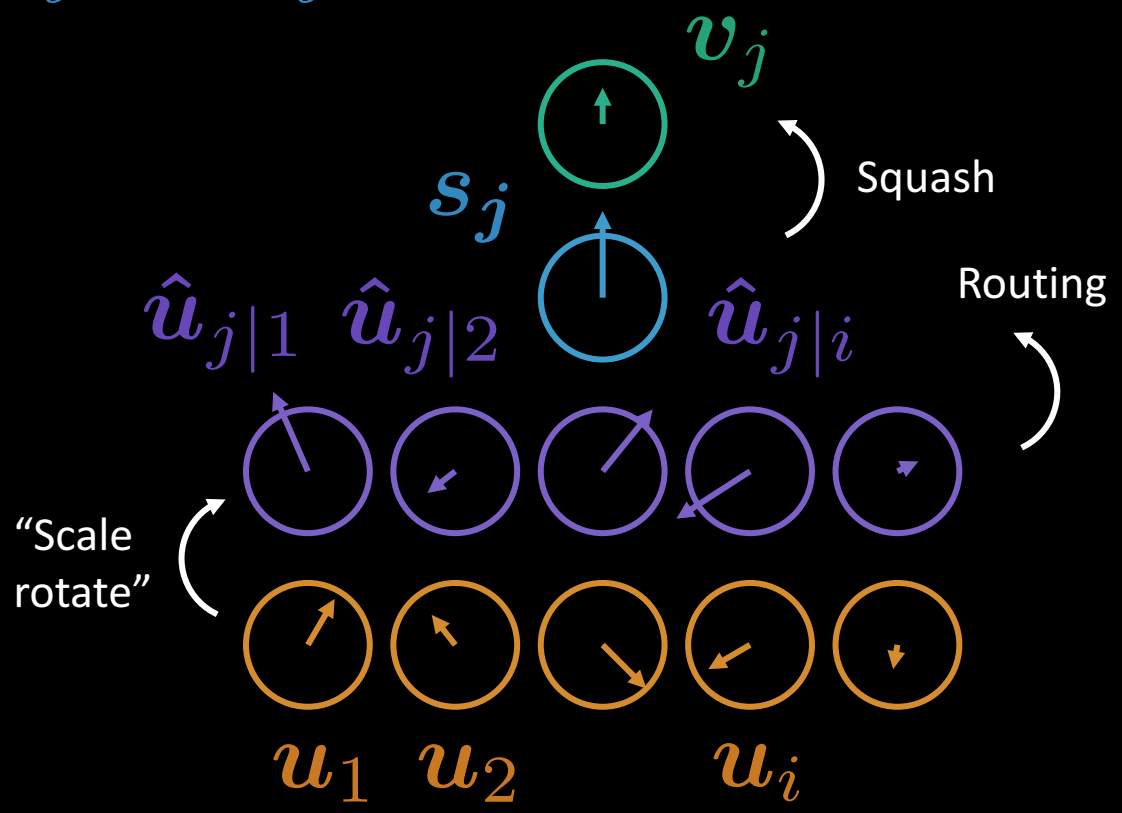

Dynamic routing

$$\boldsymbol{b}_i \leftarrow \mathbf{0}, i = 1, \ldots, s_\ell$$

$$\boldsymbol{c}_i \leftarrow \text{softmax}(\boldsymbol{b}_i)$$

$$\boldsymbol{s}_j = \sum_i c_{ij} \hat{\boldsymbol{u}}_{j|i}$$

$$\boldsymbol{v}_j \leftarrow \text{squash}(\boldsymbol{s}_j)$$

$$b_{ij} \leftarrow b_{ij} + \hat{\boldsymbol{u}}_{j|i}^\top \boldsymbol{v}_j$$

$$j = 1, \cdots, s_{\ell+1}$$

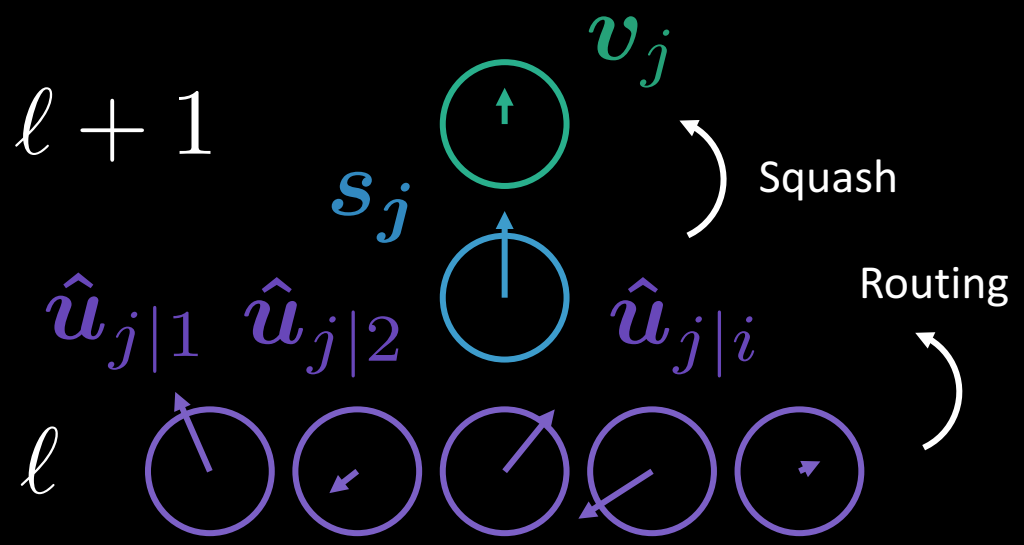

s_ℓ: size of layer ℓ

Margin loss

$$\mathcal{L}_k^{(i)} = \begin{cases} \left[\left(m_+ - \| \boldsymbol{v}_k^{(i)} \| \right)^+ \right]^2, & \text{digit } k \text{ preset} \\ \lambda \left[\left(\| \boldsymbol{v}_k^{(i)} \| - m_- \right)^+ \right]^2, & \text{otherwise} \end{cases}$$

$$\mathcal{L}_\mu^{(i)} = \sum_k \mathcal{L}_k^{(i)}$$

↑ margin

CapsNet

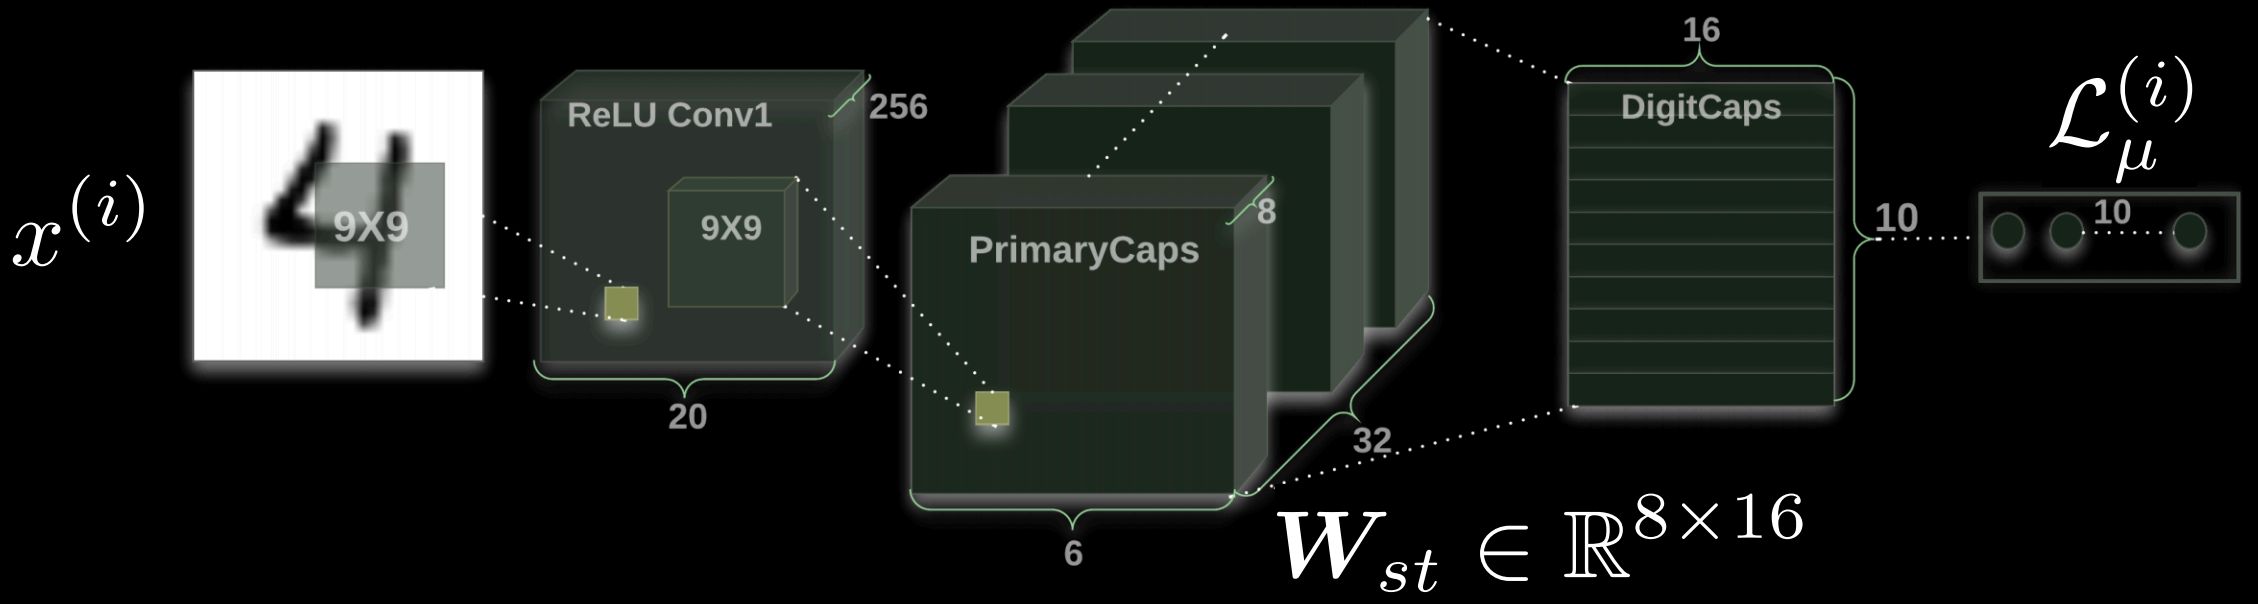

$$\boldsymbol{W}_{st} \in \mathbb{R}^{8\times 16}$$

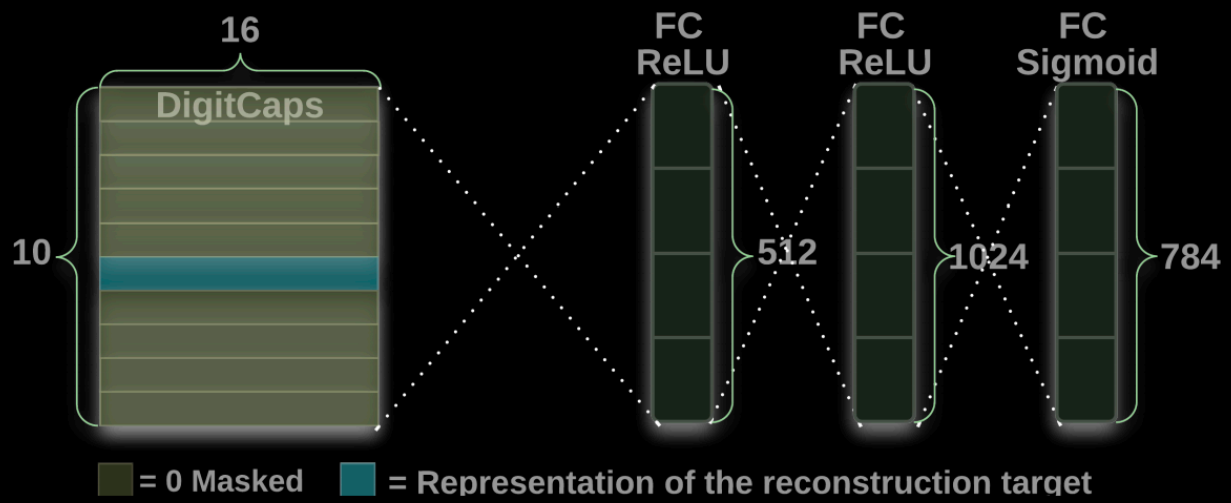

$$\mathcal{L}^{(i)} = \mathcal{L}_{\mu}^{(i)} + \rho \cdot \mathcal{L}_{\rho}^{(i)}$$

margin

reconstruction

Results (I)

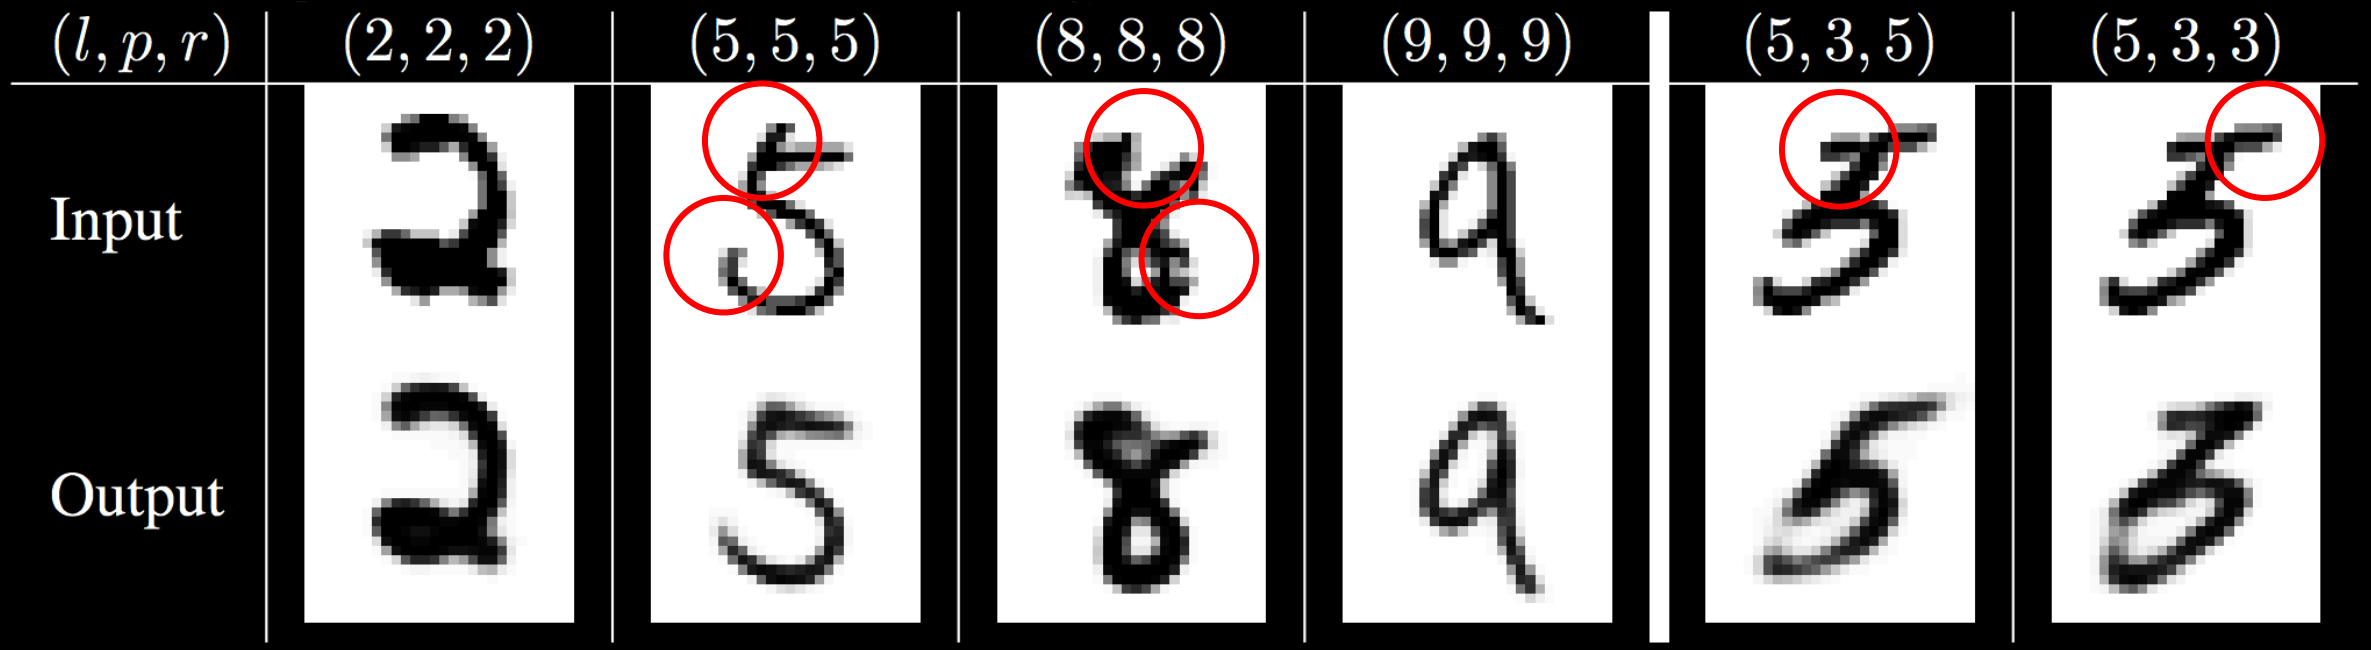

Results (II)

Scale and thickness	
Localized part	
Stroke thickness	
Localized skew	
Width and translation	
Localized part	

R:(2, 7) L:(2, 7)	R:(6, 0) L:(6, 0)	R:(6, 8) L:(6, 8)	R:(7, 1) L:(7, 1)	*R:(5, 7) L:(5, 0)	*R:(2, 3) L:(4, 3)	R:(2, 8) L:(2, 8)	R:P:(2, 7) L:(2, 8)
							
R:(8, 7) L:(8, 7)	R:(9, 4) L:(9, 4)	R:(9, 5) L:(9, 5)	R:(8, 4) L:(8, 4)	*R:(0, 8) L:(1, 8)	*R:(1, 6) L:(7, 6)	R:(4, 9) L:(4, 9)	R:P:(4, 0) L:(4, 9)